

Porting your Pytorch application to Voyager Gaudi Architecture

Javier Hernandez-Nicolau
Computational Data Scientist

Advanced HPC-CI Webinar Series. April 8th 2025

Outline

- Introduction to the Voyager system
- Containers and Kubernetes
- Running a pod on Voyager
- Porting a Pytorch application
- Parallel runs with MPI
- Jupyter Notebooks
- Import a model from Huggingface

Voyager: an HPC system with an innovative architecture dedicated to AI

- Envisioned as a system to facilitate exploration of a **new architectures** in support of **AI** in research and engineering.
- Its main component is the **Intel Gaudi** accelerators. They are specifically **designed** for **AI** applications and its optimized software allows users accelerate and scale their models.
- Voyager is a NSF category II HPC system. Test phase + production.
- During the 3-year **test phase**, the Voyager team has engaged with the research community. A **large number** of Deep Learning **models** from multiple fields have been deployed and analyzed on Voyager allowing us to learn more about this unique architecture.
- In the **allocation** (production) phase, Voyager will open its doors to a broader community.



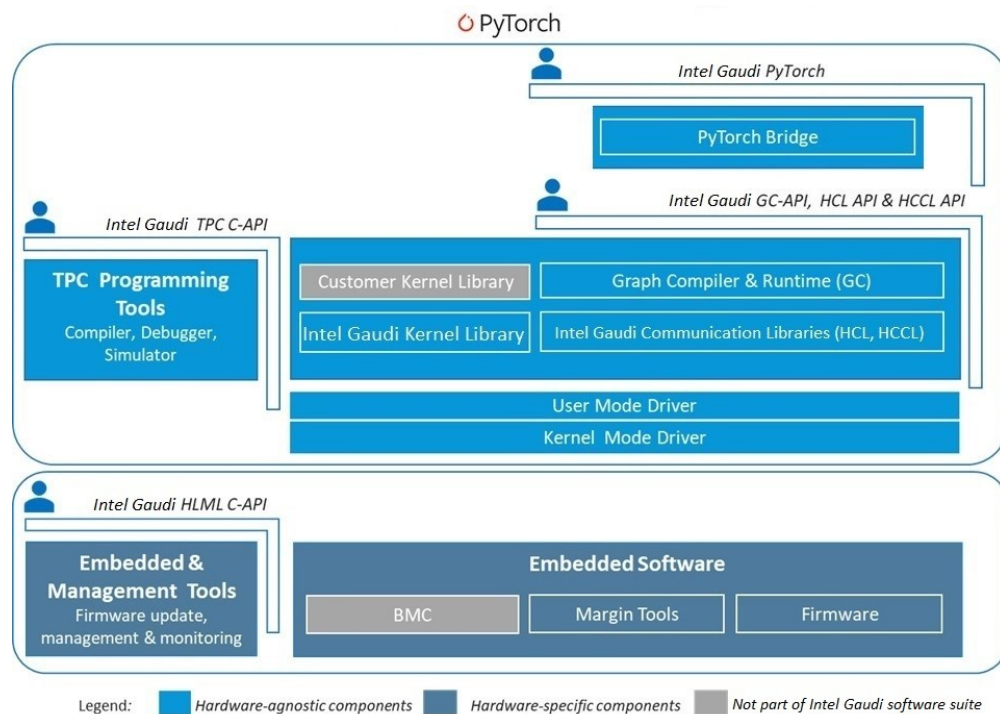
Voyager system: architecture

- 42 **Training nodes**, each with 8 **Intel Gaudi** cards specialized for AI. All-to-all networks between processors in a node.
- 36 Intel x86 processors **compute nodes** for general purpose (e.g. data processing...)
- 400 Gbe interconnect using RDMA over converged **Ethernet**
- Storage:
 - Gaudi nodes: 6.4Tb **local NVME**
 - **Ceph**: with 3PB connected via 25Gbe
 - **Home**: 324 TB connected to compute via 25Gbe

System component	Configuration
<i>Supermicro X12 Gaudi Training Nodes</i>	
CPU Type	Intel Xeon Gold 6336
Intel Gaudi processors	336
Nodes	42
Training processors/Node	8
Host X86 processors/node	2
Memory capacity	512 GB DDR4 DRAM
Memory/training processor	32 GB HDM2
Local storage	6.4 TB local NVMe
Compute nodes	
CPU Type	Intelx86
Nodes	36
X86 preocessors/node	2
Memory capacity	384 GB
Local NVMe	3.2 TB
Storage System	
Ceph	3PB
Home	324 TB

Software stack overview: Synapse AI

- At a very high level, users need very **little code modification** to run their AI models.
- **Pytorch** is **natively** integrated to Synapse AI and allows to run many popular AI applications in computer vision and generative AI (LLMs, Diffusion models).
- **Advanced users** can write their own customized kernel.
- Synapse AI is periodically **updated** including more features.



<https://docs.habana.ai/>

Scientific applications already running on (or being ported to) Voyager

Field/Project	AI architecture
Cosmology: Super resolution of galactic simulations	Diffusion
Biomedical text analytics	LLM
Cardiac image analysis	U-Net
High energy physics	Graph NN
Molecular simulations	NN
Data driven weather prediction	U-Net
Human microbiome research	Categorical VAE
Dose prediction in cervical brachytherapy	U-Net
Ontology-driven finetuning of LLM for epilepsy	LLM
Research accessibility via visual representation	Diffusion

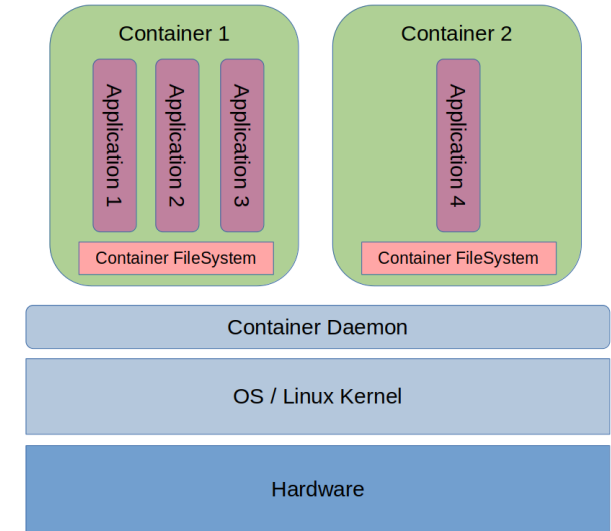
How to access Voyager

- Feel free to contact us at consult@sdsc.edu
- **Allocation** time can be requested via NAIRR. Several projects are already running.
- Starting June 1st, Voyager will get in to **production phase**. Allocations can be requested via ACCESS.
- Voyager can be accessed with **ssh**:
 - ssh to login.voyager.sdsc.edu
 - On Windows use any ssh app. For linux/Mac use the terminal
- **Tutorials** and reference models can be found at:
<https://github.com/javierhndev/Voyager-Reference-Models>

Containers and Kubernetes

What is a container?

- A container is a way to **package applications** with all the necessary dependencies and configurations.
- In a container you can install your **own** software/libraries or specific versions.
- They are **portable** and can be easily shared with others. Also containers simplify **reproducibility**.
- Is this a Virtual Machine? No!
 - **VM**: Virtualization at hardware level. Isolation of machines
 - **Containers**: Virtualization at OS level. Isolation of processes.



Kubernetes

- Containers on Voyager are deployed via **Kubernetes** (K8S).
- Kubernetes is an **orchestration** system to deploy and manage containers.
- Kubernetes **objects** describe the containerized applications, software, resources (memory, cpus...), volumes, policies....
- K8S objects are described by a **configuration file** in json or yaml format.
- The smallest object in K8S is a **pod** and includes one (or more) containers.
- Each container must pick an **image** which includes all the **necessary software** to run the application.
- Images are **externally hosted**. The most famous repository being DockerHub. However images for Voyager and Intel Gaudi can be obtained from vault.habana.ai.



kubernetes

Manage applications in Kubernetes

- The `kubectl` command can be used to create, **manage** and delete kubernetes objects.
- It can be loaded on Voyager with the `module load kubernetes/voyager` command
- To **create** an object:
 - `kubectl create -f mypod.yaml`
- To **delete** the object:
 - `kubectl delete -f mypod.yaml`
- **List** all pods:
 - `kubectl get pods`
- Get the pod **logs**:
 - `kubectl logs podname`



Hello world pod

```
apiVersion: v1
kind: Pod
metadata:
  name: test
spec:
  restartPolicy: Never
  containers:
    - name: gaudi-container
      image: vault.habana.ai/audi-docker/1.15.1/ubuntu22.04/habanalabs/pytorch-
installer-2.2.0:latest
      resources:
        limits:
          memory: 32G
          cpu: 12
          habana.ai/audi: 1
      command: ["/bin/sh", "-c"]
      args:
        - echo 'Hello World!';
```

Pod name

What if pod crashes?

Container image to use

Gaudi cards

Commands to be executed

Run the hello world pod

- Launch a single pod with:
 - `Kubectl create -f podfile.yaml`
- Check their status:
 - `Kubectl get pods`

NAME	READY	STATUS	RESTARTS	AGE
test	0/1	Completed	0	57s

- Check the output log with:
 - `kubectl logs podname`

A more complex yaml file

```
apiVersion: v1
kind: Pod
metadata:
  name: mnist
spec:
  restartPolicy: Never
  volumes:
    - name: scratch
      emptyDir: {}
    - name: workdir
      hostPath:
        path: /home/javierhn/webinar_adv_comp_2025
        type: Directory
  containers:
    - name: gaudi-container
      image: vault.habana.ai/audi-docker/1.15.1/ubuntu22.04/habanalabs/pytorch-
installer-2.2.0:latest
    ....
    .... (continue next page)
```

Diagram illustrating the structure of the YAML file and its components:

- Folders container can access** (points to `restartPolicy: Never`)
- Scratch (purged)** (points to `name: scratch`)
- Folder from /home FS** (points to `path: /home/javierhn/webinar_adv_comp_2025`)

A more complex yaml file (continue)

containers:

- name: gaudi-container

- image: vault.habana.ai/audi-docker/1.15.1/ubuntu22.04/habanalabs/pytorch-installer-2.2.0:latest

volumeMounts:

- mountPath: /scratch
name: scratch
- mountPath: /workdir
name: workdir

Mounts folders
in the container

Defines name
for folders

resources:

limits:

memory: 32G
cpu: 12
habana.ai/audi: 1
hugepages-2Mi: 95000Mi

requests:

memory: 32G
cpu: 12
habana.ai/audi: 1
hugepages-2Mi: 95000Mi

Gaudi card info

command: ["/bin/sh", "-c"]

args:

- hl-smi;
- cd /workdir;
- #pip install apackage;
- python run_mnist_onecard.py --epochs 20 --target-accuracy 0.90;

Install
dependencies

Porting a Pytorch application to the Intel Gaudi architecture

Porting a Pytorch application to Intel Gaudi architecture

- Import the Intel Gaudi Pytorch framework:
 - `import habana_frameworks.torch.core as htcore`
- Target the Gaudi device:
 - `device = torch.device("hpu")`
- A `mark_step()` must be added after every `loss.backward()` and `optimizer.step()` during training and after loss calculation when testing/validating:
 - `htcore.mark_step()`
- See example (CNN with MNIST dataset) in repo.

Porting a Pytorch application to Intel Gaudi architecture

- More **complicated** model? CUDA/GPU API calls?
 - GPU Migration Toolkit
- **Deepspeed?**
 - Use Intel Gaudi's forked version
- Distributed training (DDP) with **NCCL**?
 - Use HCCL equivalent
- You can always contact SDSC support via consult@sdsc.edu

MPI parallelization on Voyager

Distributed Parallel training with MPI on Voyager

- Voyager is able to **scale up** your application to multiple Gaudi cards (and nodes!).
- For distributed training on (or less than) 8 cards you can use the Kubernetes **pod** object with the `mpirun` command.
- Running in **multiple nodes** (16 cards or more) requires a Kubeflow MPIJob.
- Check out our tutorials on <https://github.com/javierhndev/Voyager-Reference-Models>

Jupyter notebooks

Jupyter notebooks on Voyager

- Launch a **pod** with Jupyter:

- `pip install jupyter;`
- `jupyter notebook --allow-root;`

To access the server, open this file in a browser:

`file:///root/.local/share/jupyter/runtime/jpserver-48-open.html`

Or copy and paste one of these URLs:

`http://localhost:8888/tree?token=87c14bf9148425dcc3a1dc5d6cd5b9b1`

`http://127.0.0.1:8888/tree?token=87c14bf9148425dcc3a1dc5d6cd5b9b1`

- **Check** the Jupyter links with

- `kubectl logs yourpod`

- Port forward the pod:

- `kubectl port-forward podname 9888:8888`

- Create an **ssh tunnel** from the client (your pc):

- `ssh -N -f -L 8888:localhost:9888 login.voyager.sdsc.edu`

- Copy the Jupyter link to **your browser**.




Huggingface models on Voyager

Optimum-habana



Hugging Face

- **LLMs** and **Diffusion** models from **Huggingface** can be run on Voyager with the **optimum-habana** package.
- Optimum-habana (<https://github.com/huggingface/optimum-habana>) is built on top of **Huggingface Transformers/diffusers** and allows the models to take advantage of the **Gaudi** accelerators.



- **LLM:** finetuning of Llama3:8B, Llama3.1:8B, Gemma2:27B, Deepseek-R1-DistillQwen-32B
- Text Generation Inference (TGI)
- **Diffusers:** Stable Diffusion 1.5, Stable Diffusion XL1, CommonCanvas-XL-C
- **LMM:** qwen2-VL-7B, llama3.2-11B, llava-1.6-Mixtral-7B

Optimum-habana



Hugging Face

```
from optimum.habana.diffusers import GaudiDDIMScheduler, GaudiStableDiffusionPipeline

model_name='/dataset/stable-diffusion-v1-5'

scheduler = GaudiDDIMScheduler.from_pretrained(model_name, subfolder="scheduler")

pipeline = GaudiStableDiffusionPipeline.from_pretrained(
    model_name,
    scheduler=scheduler,
    use_habana=True,
    use_hpu_graphs=True,
    gaudi_config="Habana/stable-diffusion",
)

outputs = pipeline(
    ["A panda eating a taco"],
    num_images_per_prompt=8,
    batch_size=4,
)

image_save_dir="."
for i, image in enumerate(outputs.images):
    image.save( f"image_{i+1}.png")
```

Useful links

- Intel Gaudi documentation:
 - <https://docs.habana.ai/>
- Pytorch on Intel Gaudi:
 - <https://docs.habana.ai/en/latest/PyTorch/index.html>
- Reference models/Tutorials:
 - Intel-Gaudi: <https://github.com/HabanaAI/Model-References>
 - Voyager: <https://github.com/javierhndev/Voyager-Reference-Models>
- Optimum-habana:
 - <https://github.com/huggingface/optimum-habana>

Summary

- **Intel Gaudi** accelerators, designed for **AI**, are the main component of Voyager.
- Applications in Voyager are run in **containers** and are managed by **Kubernetes**. `kubectl` is the main command.
- Users can port their **Pytorch** applications with **little effort**.
- **Jupyter** notebook is **available** on Voyager
- To run LLMs or Diffusers from **Huggingface**, users need **optimum-habana** for Intel Gaudi.