



Advanced HPC-CI Webinar

Scalable Machine Learning - Spark

Mai H. Nguyen, Ph.D.

2025-01

Spark Introduction

Spark Topics

- **Spark**
 - History
 - RDDs
 - DataFrames
 - Spark Design Goals
 - Spark API
 - Spark Core & Libraries
- **Spark Demo**
 - Scaling
 - Cluster Analysis

- **Computing platform for scalable computing**
 - Designed for big data workloads
 - Built-in parallelism & fault-tolerance on commodity cluster
 - Provides interactive querying, iterative analytics, streaming processing, along with batch processing
 - Goals: speed, ease of use, generality, unified platform
- **History**
 - Research project began in 2009 at UC Berkeley's AMPLab
 - Paper published in 2010
 - Contributed to Apache Software Foundation in 2013
 - Commercial version by Databricks

SPARK

- Goals: **speed**, ease of use, generality, unified platform
- In-memory processing
 - Exploits distributed memory to cache data
 - Intermediate results written to memory whenever possible
- How does Spark manage data in distributed system?

RESILIENT DISTRIBUTED DATASETS (RDDs)

- Spark central concept
 - Abstraction of data as distributed collection of objects
- Resilient Distributed Datasets (RDDs)
 - Data abstraction
 - Programming construct for storing and organizing data
 - Spark uses RDDs to distribute data and computations across nodes in cluster

RDD

- Resilient Distributed **Dataset**
 - Collection of data
 - From files in local filesystem (text, JSON, etc.)
 - From data store (HDFS, RDBMS, NoSQL, etc.)
 - Created from another RDD
- Resilient **Distributed** Dataset
 - Data is divided into partitions
 - Partitions are distributed across nodes in cluster
- **Resilient** Distributed Dataset
 - Provides resilience (e.g., fault tolerance) to failures
 - History of operations performed on each partition is tracked to provide lineage-based fault tolerance
- All provided automatically by Spark engine

DATAFRAMES & DATASETS

- **Extensions to RDDs**
 - Higher-level abstractions
 - Improved performance
 - Better scalability
- **Can convert to/from RDDs and use with RDDs**

DATAFRAMES & DATASETS

DataFrame

- Lazy evaluation
- Immutable
- Data organized as collection of Rows
- No static type checking
- APIs in Java, Scala, Python, R

DataSet

- Lazy evaluation
- Immutable
- Data organized as collection of Rows
- Provides static type checking
- APIs in Java and Scala

USING DATAFRAMES

- Spark Session
 - Entry point to Spark engine
 - Note that SparkContext is now **SparkSession**

```
from pyspark import SparkSession, SparkConf

conf = SparkConf \
    .setAll \
    ([("spark.app.name", "DataFrame Example") \
      ("config.option", "config.value")])

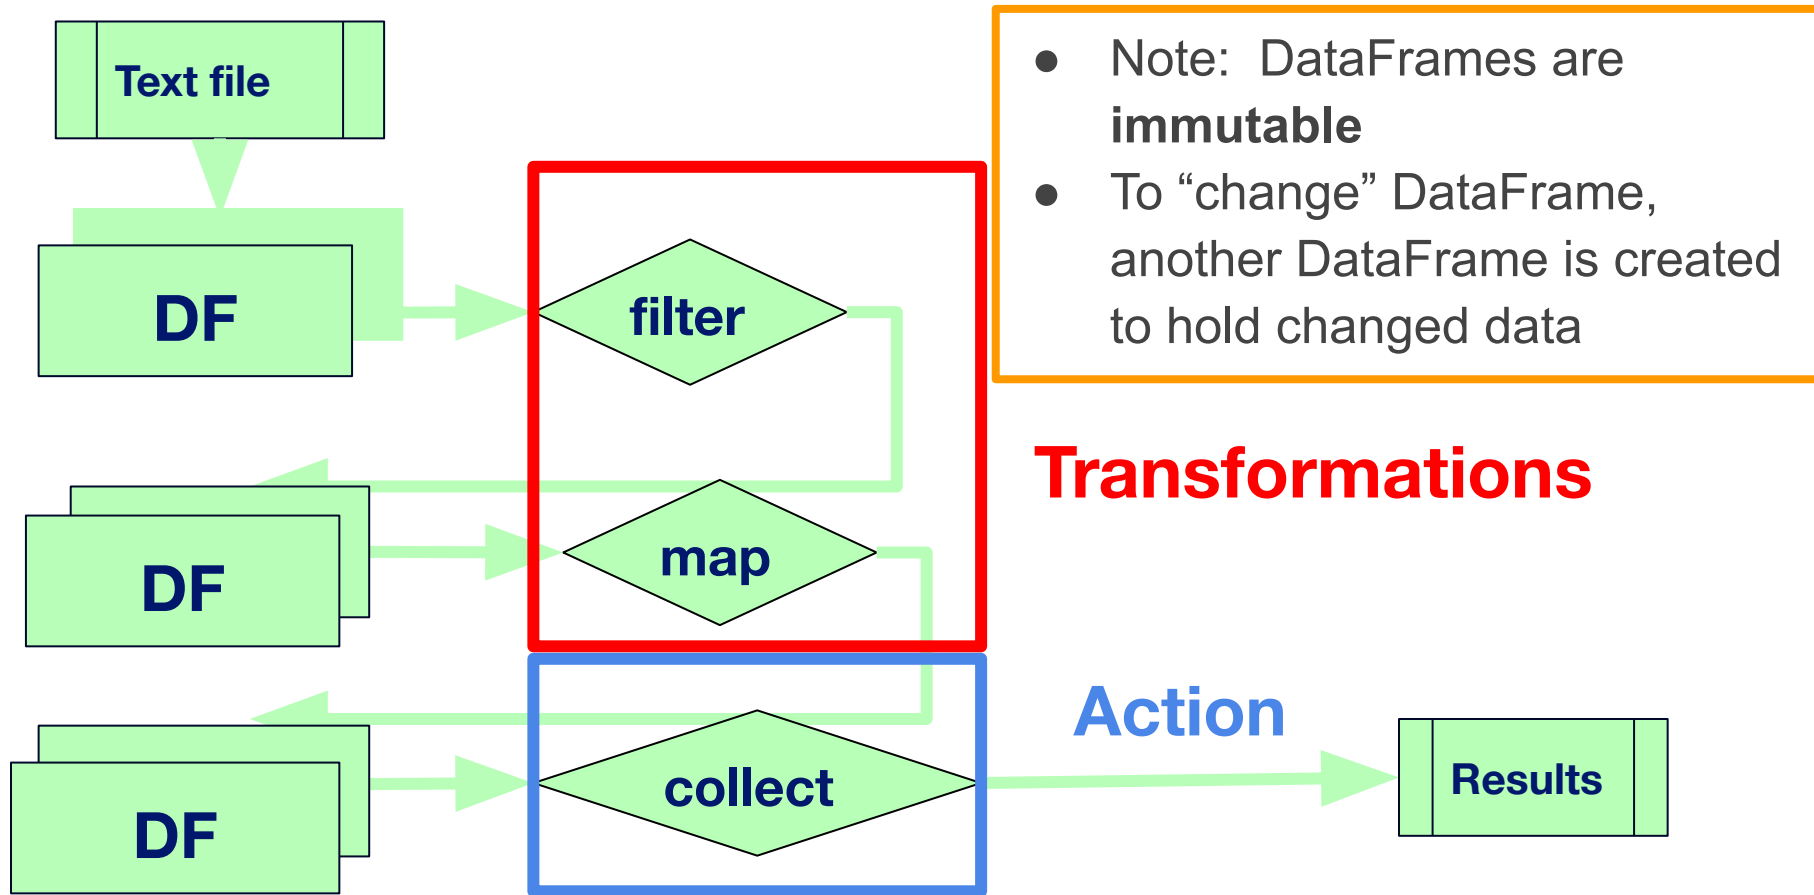
spark =
    SparkSession.builder.config(conf=conf) \
        .getOrCreate()
```

CREATING DATAFRAMES

- Read data from files in local filesystem (text, JSON, etc.)
 - `df = spark.read.csv("data.csv", header="True")`
- Data read in from data store (HDFS, RDBMS, NoSQL, etc.)
 - `df = spark.read.csv("hdfs:///<path>/data.csv")`
- Generate data
 - `empl_0 = Row(id="123", name="John")`
 - `empl_1 = Row(id="456", name="Mary")`
 - `employees = [empl_0, empl_1]`
 - `df = spark.createDataFrame(employees)`
- Created by transforming another DataFrame
 - `filter_df = df.filter(col("name")== "Mary")`

PROCESSING DATAFRAMES

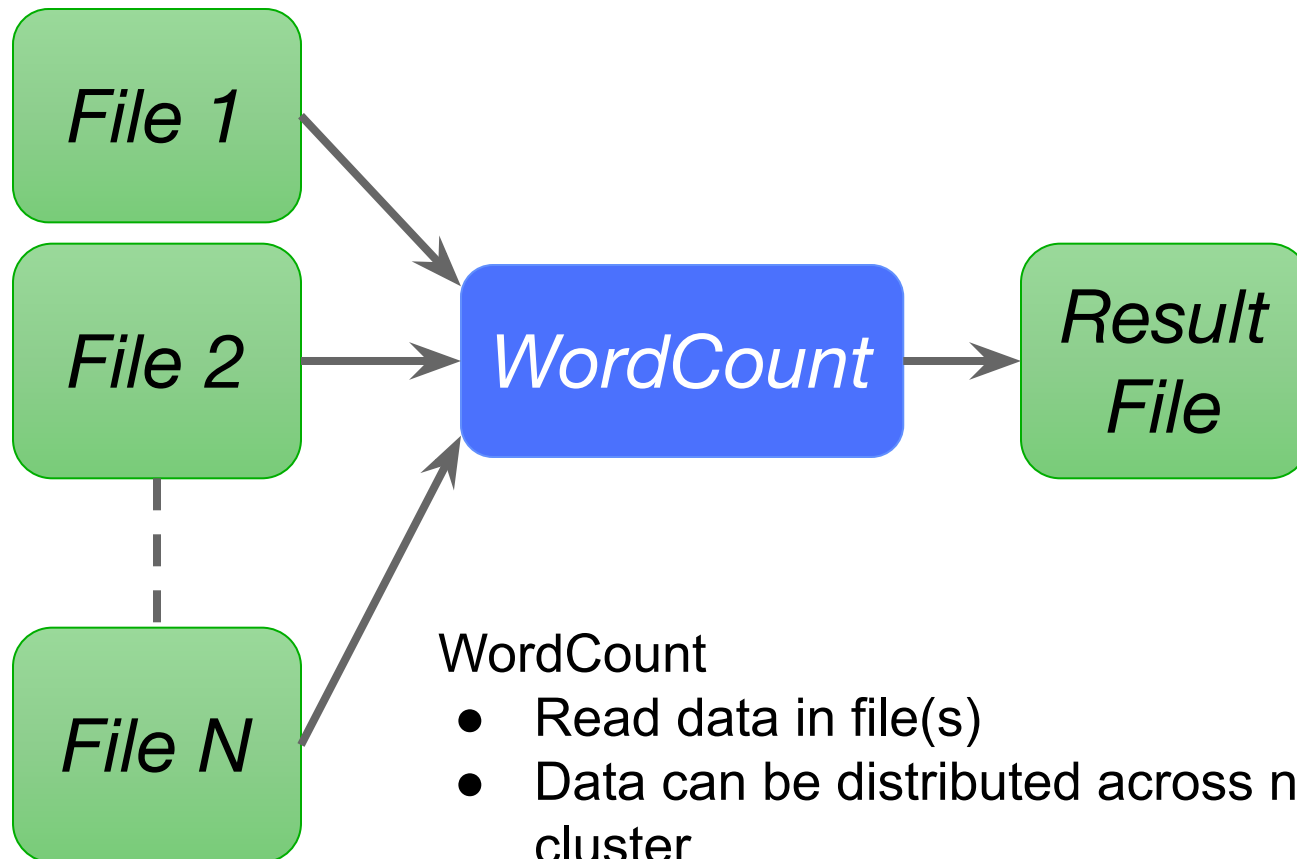
- DataFrames can be processed using 2 types of operations
 - **Transformation:** Creates new DataFrame from existing DataFrame
 - **Action:** Runs computation(s) on DataFrame and returns value



LAZY EVALUATION

- Transformations on DataFrames have **lazy evaluation**
 - Transformations are not immediately processed
 - Plan of operations is built
- Operations executed when **action** is performed
 - i.e., actions force computation
- Allows for optimizations in generating physical plan
- Example:
 - `filtered = strings.filter(strings["value"].contains("Spark"))`
 - Nothing is returned
 - `filtered.count()`
 - 'filter' is performed, and count is returned

WordCount

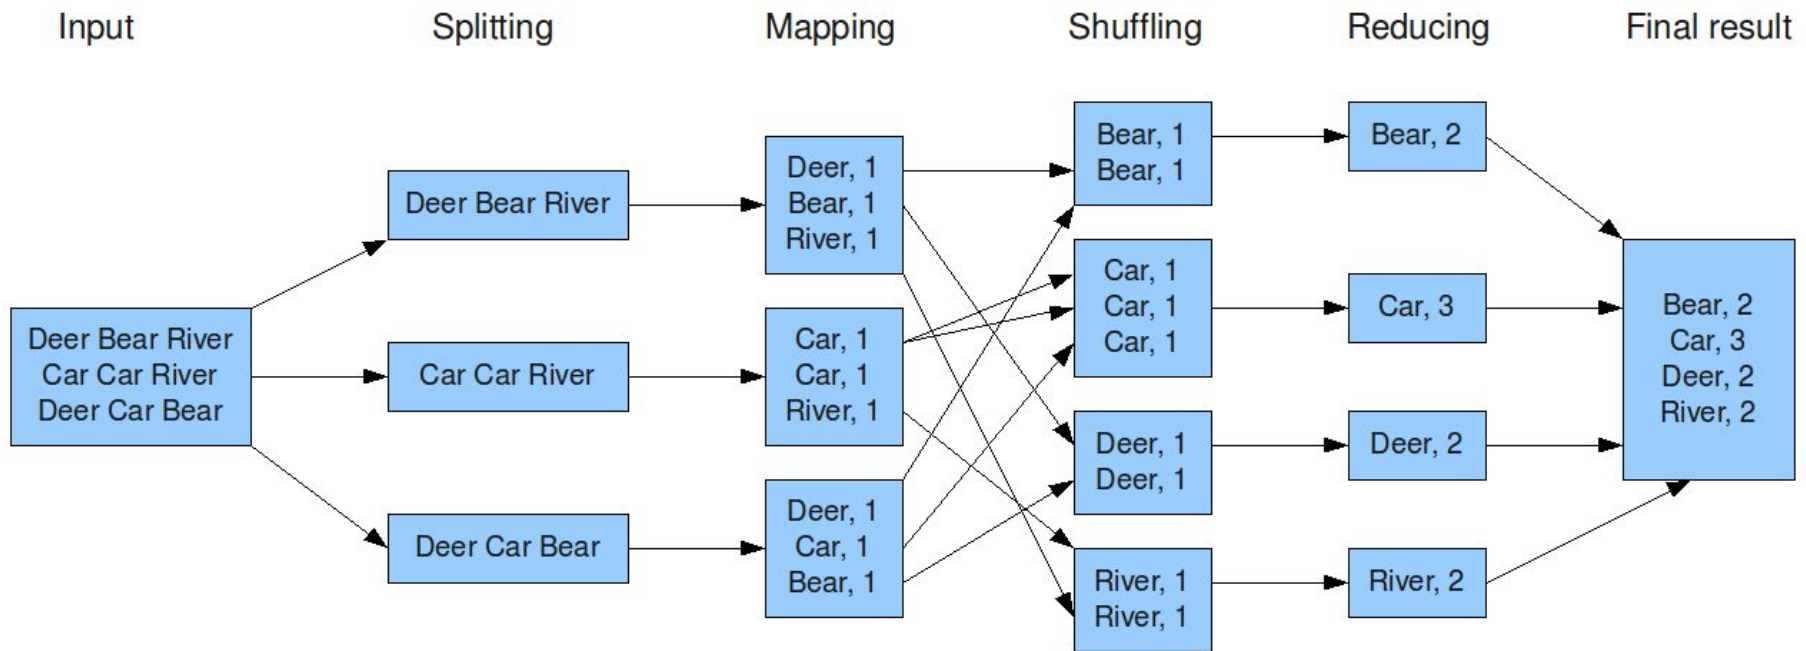


WordCount

- Read data in file(s)
- Data can be distributed across nodes in cluster
- Count number of occurrences of each word

WordCount

The overall MapReduce word count process



<https://www.todaysoftmag.com/article/1358/hadoop-mapreduce-deep-diving-and-tuning>

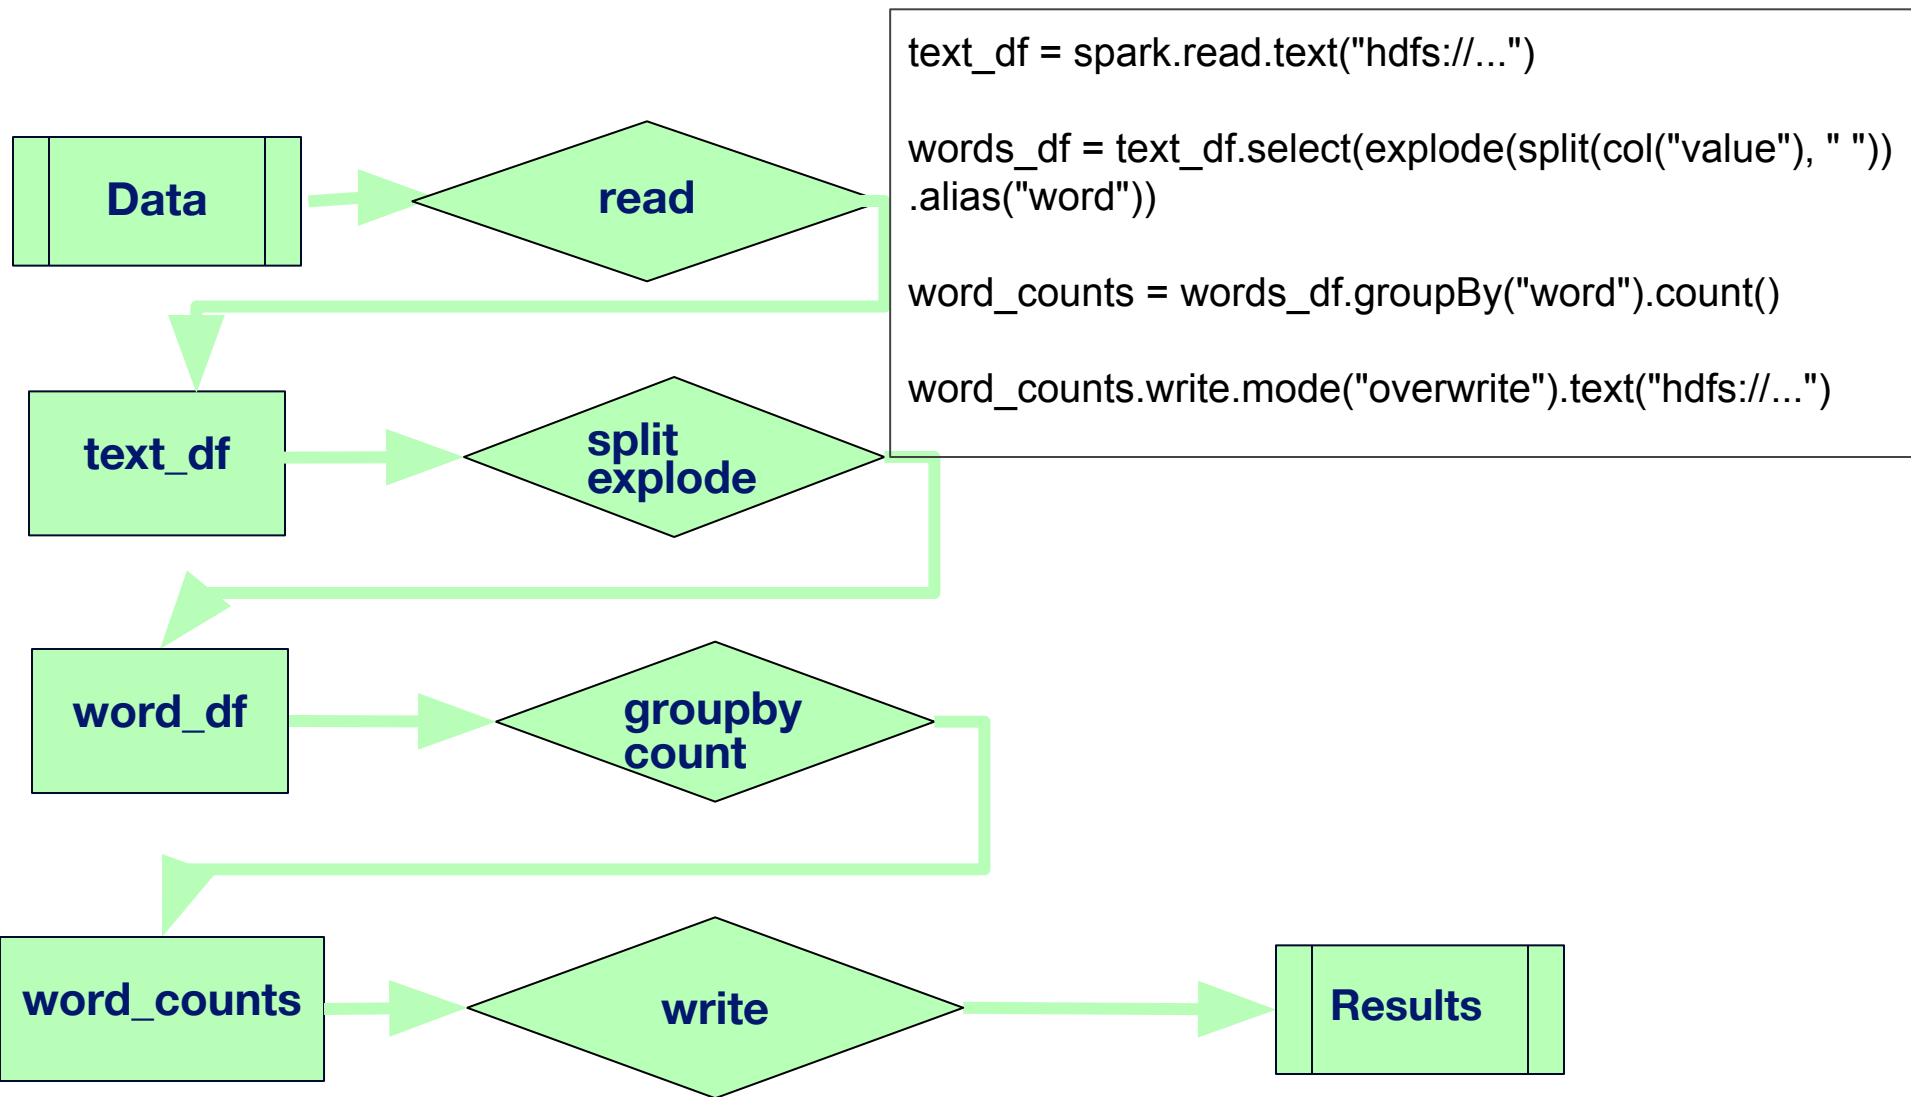
Data is split into
partitions

Map generates
key-value pairs

Pairs with same
key moved to
same partition

Reduce sums
values for
each key


WordCount



SPARK PROGRAM STRUCTURE


- **Start Spark session**
 - `spark = SparkSession.builder.config(conf=conf).getOrCreate()`
- **Create distributed dataset**
 - `df = spark.read.csv("data.csv",header="True")`
- **Apply transformations**
 - `new_df = df.filter(col("dept") == "Sales")`
- **Perform actions**
 - `df.collect()`
- **Stop Spark session**
 - `spark.stop()`

SPARK ARCHITECTURE



SPARK INTERFACE

Goals: speed, **ease of use**, generality, unified platform



WORDCOUNT EXAMPLE IN SPARK

Spark API available in Python, Scala, Java, and R

PySpark

```
text_df = spark.read.text("hdfs://...")
words_df = text_df.select(explode(split(col("value"), " ")).alias("word"))
word_counts = words_df.groupBy("word").count()
word_counts.write.mode("overwrite").text("hdfs://...")
```

SparkR

```
textDF <- read.text(spark, "hdfs://...")
wordsDF <- selectExpr(textDF, "explode(split(value, ' ')) as word")
wordCounts <- count(groupBy(wordsDF, "word"))
write.df(wordCounts, "hdfs://...", "text", mode = "overwrite")
```

Scala


```
val textDF = spark.read.text("hdfs://...")
val wordsDF = textDF.select(explode(split(col("value"), ""))).alias("word"))
val wordCounts = wordsDF.groupBy("word").count()
wordCounts.write.mode("overwrite").text("hdfs://...")
```

SPARK - GENERALITY

- Goals: speed, ease of use, **generality**, unified platform
- Support for several data sources
 - Local file systems, HDFS, RDBMSs, MongoDB, Kafka, AWS S3, etc.
- Can run on various platforms
 - Hadoop, Kubernetes, cloud, standalone
- Support for multiple workloads
 - batch, streaming
 - machine learning, SQL, graph processing


SPARK - UNIFIED PLATFORM

- Goals: speed, ease of use, generality, **unified platform**



- Provides unified platform for various analytics processing
- **Spark engine** provides core capabilities for scalable processing
- **Spark libraries** provide additional higher-level functionality for diverse workloads

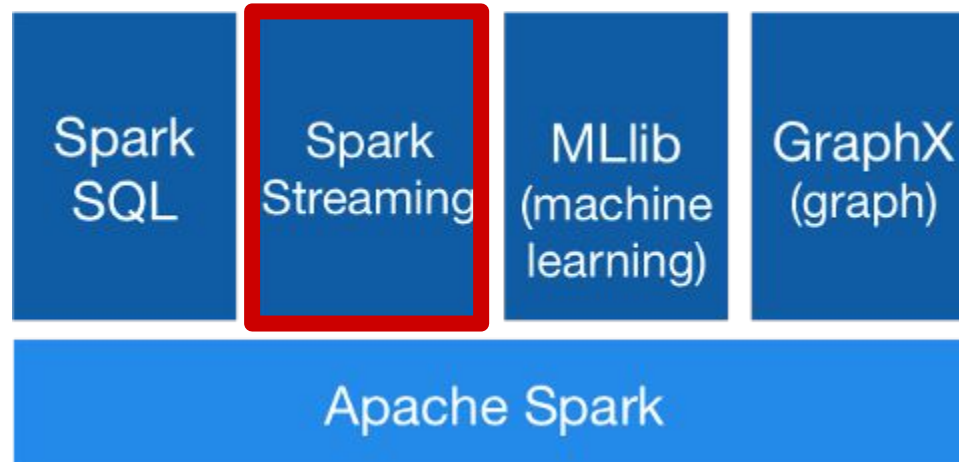
SPARK SQL



- **Structured Data Processing**

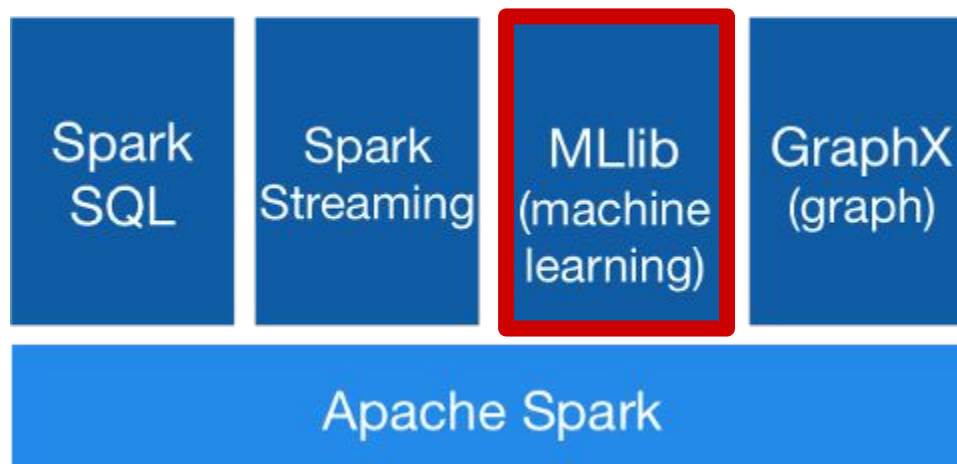
- Provides support for SQL and query processing
- Has APIs for SQL, Scala, Java, Python, and R
- Generated underlying code is identical

SPARK STREAMING



- **Streaming Data Processing**
 - Scalable processing for real-time analytics
 - Structured streaming
 - Data stream is divided into micro-batches of data
 - Same operations for static data can be used
 - Has APIs for Scala, Java, and Python

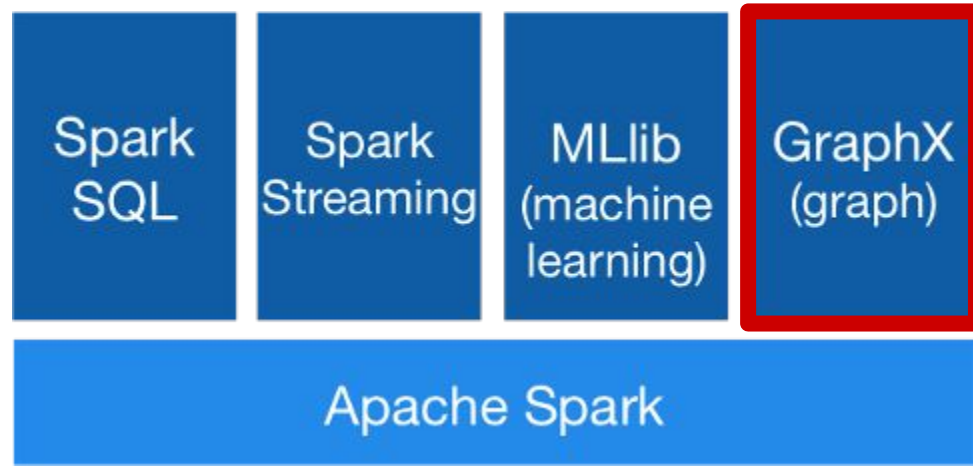
SPARK MLLIB



- **Machine Learning**

- Scalable machine learning library
- Scalable implementations of machine learning algorithms and utilities
- Has APIs for Scala, Java, Python, and R

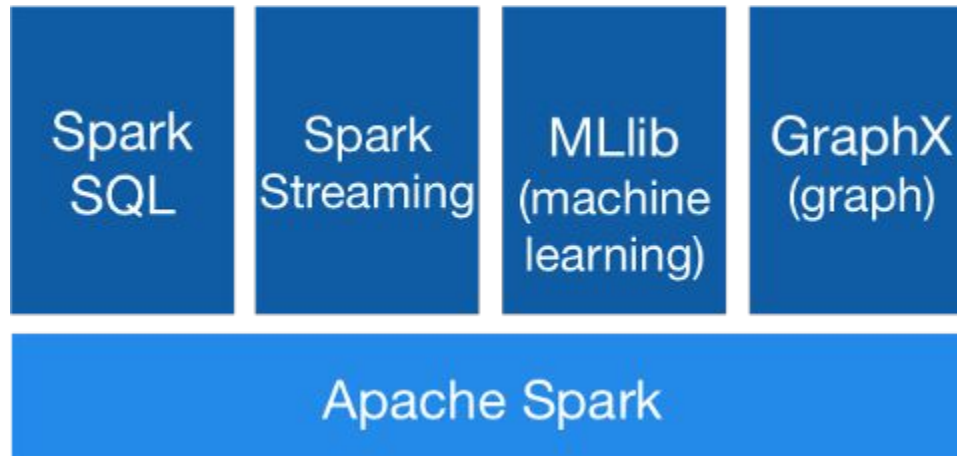
SPARK GRAPHX / GRAPHFRAMES



- **Graph Computation**

- Scalable graph processing
- Special structures for storing vertex and edge information & operations for manipulating graphs
- GraphX (RDD-based) & GraphFrames (DF-based)
- Has APIs in Scala, Java, Python (GraphFrames)

SPARK



Unified engine for large-scale data analytics

Goals: speed, ease of use, generality, unified platform

Spark Resources

- PySpark SQL Basics Cheat Sheet
 - PDF
- Spark Main Page
 - <https://spark.apache.org/>
- Spark Overview
 - <https://spark.apache.org/docs/latest/index.html>
- Spark Examples
 - <https://spark.apache.org/examples.html>
- Spark SQL, DataFrames and DataSets Programming Guide
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- Spark MLlib Programming Guide
 - <https://spark.apache.org/docs/latest/ml-guide.html>
- PySpark API Documentation
 - <https://spark.apache.org/docs/latest/api/python/index.html>

Spark Demo

Server Setup for PySpark - Command Line

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu -l <account>`

- **In terminal window**

- `export PATH="/cm/shared/apps/sdsc/galileo:${PATH}"`
- `jupyter-shared-spark`
 - Alias for: `galileo launch --account ${HPC_ACCOUNT} --reservation ${HPC_RESERVATION_CPU} --partition shared --cpus 4 --memory 16 --time-limit 02:00:00 --env-modules singularitypro --sif /cm/shared/apps/containers/singularity/spark/spark-latest.sif --bind /expanse,/scratch,/cm --quiet`

- **To check queue**

- `squeue -u $USER`

PySpark Scaling Hands-On

- **Data**
 - Book reviews
 - Source : <https://jmcauley.ucsd.edu/data/amazon/>
- **Notebook**
 - pyspark_demo_soln.ipynb
- **To do**
 - Change number of cores: 1, 2, 4
 - Note difference in execution times
 - Run each configuration 3 times

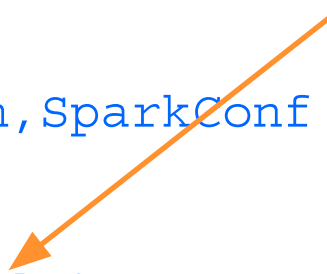
SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession, SparkConf

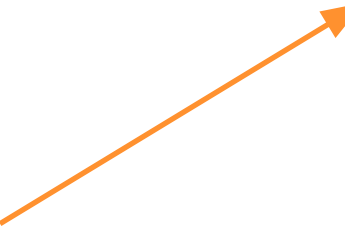
conf = SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])

spark = SparkSession.builder.config(conf=conf).getOrCreate()
```


Use * to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one



GETTING EXECUTION TIMES

- In notebook, execution time is printed out in cell before Spark session is stopped (next to last cell)
- Need to restart the kernel and run all cells without stopping to get accurate execution time:
 - Run -> Restart Kernel and Run All Cells
- Find mean and standard deviation of execution times over 3 runs for
 - 1 core, 2 cores, and 4 cores

```
import pyspark
from pyspark.sql import SparkSession
```

```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[2]'),
    ('spark.app.name', 'PySpark Demo)])
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

Specify number of cores.
“*” uses all available cores



PySpark Cluster Analysis Hands-On

- **Data**
 - Weather station measurements
- **Task**
 - Perform cluster analysis to identify different weather patterns
- **Approach**
 - Spark k-means
- **Notebooks**
 - `pyspark-clustering-soln.ipynb`

Dataset Description

- **Measurements from weather station on Mt. Woodson, San Diego**
- **Air temperature, humidity, wind speed, wind direction, etc.**
- **Three years of data: Sep. 2011 - Sep. 2014**
 - measurements every minute
- **Source**
 - <http://hpwren.ucsd.edu>

A topographic map of the San Diego region, showing various peaks and locations. A red arrow points to Mt. Woodson, and a black arrow points to San Clemente Island. Other labeled locations include Wildomar, Murrieta, Temecula, SMER North Station, Red Mt., Fallbrook, Boucher Hill, Puerta La Cruz, Sky Oaks, Warner Springs, Mesa Grande, Big Black Mt., Santa Ysabel, Julian, Mt. Laguna, Lyons Peak, Alpine, Pine Valley, El Cajon, La Presa, Lemon Grove, San Diego, Encinitas, Carlsbad, Vista, Escondido, San Marcos, Valley Center, Ramona, Poway, Santee, and San Clemente Island. The map also shows the Colorado Desert, San Jacinto Mountains, and various reservations like the Colorado River Indian Reservation and the San Jacinto Mountains National Monument.

Clustering Hands-On Overview

- **Setup**
 - Start Spark
 - Load modules
- **Load data**
 - Specify schema
 - Read in data
- **Explore data**
 - Look at schema, number of rows, summary statistics
- **Prepare data**
 - Drop nulls
 - Create feature vector
- **Perform k-means cluster analysis**
 - Use elbow plot to determine k
 - Build k-means model
- **Evaluate clusters**
 - Plot cluster profiles
- **Stop Spark session**

Resources

- **Spark**
 - <https://spark.apache.org/>
- **PySpark API**
 - <https://spark.apache.org/docs/latest/api/python/index.html>
- **Spark DataFrame**
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- **MLlib**
 - <https://spark.apache.org/mllib/>
- **User's Guide**
 - https://spark.apache.org/docs/latest/api/python/user_guide/pandas_on_spark/index.html