



Advanced HPC-CI Webinar

Scalable Machine Learning - Spark

Mai H. Nguyen, Ph.D.

2025-01

Spark Introduction

Spark Topics

- **Spark**
 - History
 - RDDs
 - DataFrames
 - Spark Design Goals
 - Spark API
 - Spark Core & Libraries
- **Spark Demo**
 - Scaling
 - Cluster Analysis

- **Computing platform for scalable computing**
 - Designed for big data workloads
 - Built-in parallelism & fault-tolerance on commodity cluster
 - Provides interactive querying, iterative analytics, streaming processing, along with batch processing
 - Goals: speed, ease of use, generality, unified platform
- **History**
 - Research project began in 2009 at UC Berkeley's AMPLab
 - Paper published in 2010
 - Contributed to Apache Software Foundation in 2013
 - Commercial version by Databricks

SPARK

- Goals: **speed**, ease of use, generality, unified platform
- In-memory processing
 - Exploits distributed memory to cache data
 - Intermediate results written to memory whenever possible
- How does Spark manage data in distributed system?

RESILIENT DISTRIBUTED DATASETS (RDDs)

- Spark central concept
 - Abstraction of data as distributed collection of objects
- Resilient Distributed Datasets (RDDs)
 - Data abstraction
 - Programming construct for storing and organizing data
 - Spark uses RDDs to distribute data and computations across nodes in cluster

RDD

- **Resilient Distributed Dataset**
 - Collection of data
 - From files in local filesystem (text, JSON, etc.)
 - From data store (HDFS, RDBMS, NoSQL, etc.)
 - Created from another RDD
- **Resilient Distributed Dataset**
 - Data is divided into partitions
 - Partitions are distributed across nodes in cluster
- **Resilient Distributed Dataset**
 - Provides resilience (e.g., fault tolerance) to failures
 - History of operations performed on each partition is tracked to provide lineage-based fault tolerance
- All provided automatically by Spark engine

DATAFRAMES & DATASETS

- **Extensions to RDDs**
 - Higher-level abstractions
 - Improved performance
 - Better scalability
- **Can convert to/from RDDs and use with RDDs**

DATAFRAMES & DATASETS

DataFrame

- Lazy evaluation
- Immutable
- Data organized as collection of Rows
- No static type checking
- APIs in Java, Scala, Python, R

DataSet

- Lazy evaluation
- Immutable
- Data organized as collection of Rows
- Provides static type checking
- APIs in Java and Scala

USING DATAFRAMES

- Spark Session
 - Entry point to Spark engine
 - Note that SparkContext is now **SparkSession**

```
from pyspark import SparkSession, SparkConf

conf = SparkConf \
    .setAll \
    ([("spark.app.name", "DataFrame Example") \
      ("config.option", "config.value")])

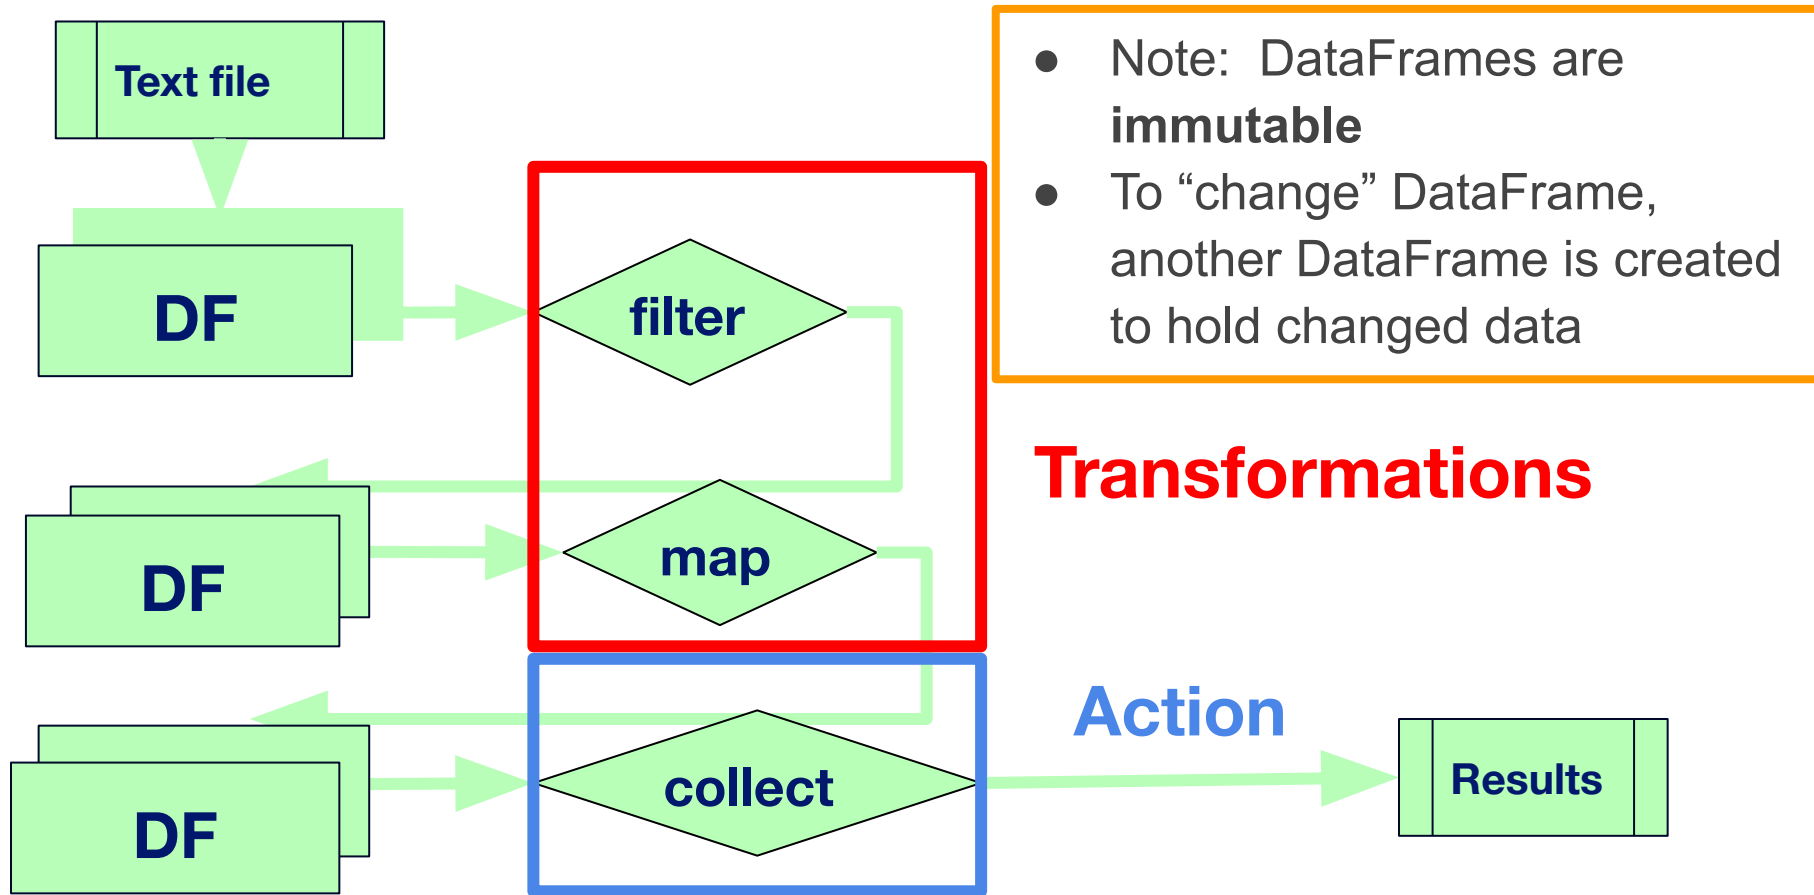
spark =
    SparkSession.builder.config(conf=conf) \
        .getOrCreate()
```

CREATING DATAFRAMES

- Read data from files in local filesystem (text, JSON, etc.)
 - `df = spark.read.csv("data.csv", header="True")`
- Data read in from data store (HDFS, RDBMS, NoSQL, etc.)
 - `df = spark.read.csv("hdfs:///<path>/data.csv")`
- Generate data
 - `empl_0 = Row(id="123", name="John")`
 - `empl_1 = Row(id="456", name="Mary")`
 - `employees = [empl_0, empl_1]`
 - `df = spark.createDataFrame(employees)`
- Created by transforming another DataFrame
 - `filter_df = df.filter(col("name")== "Mary")`

PROCESSING DATAFRAMES

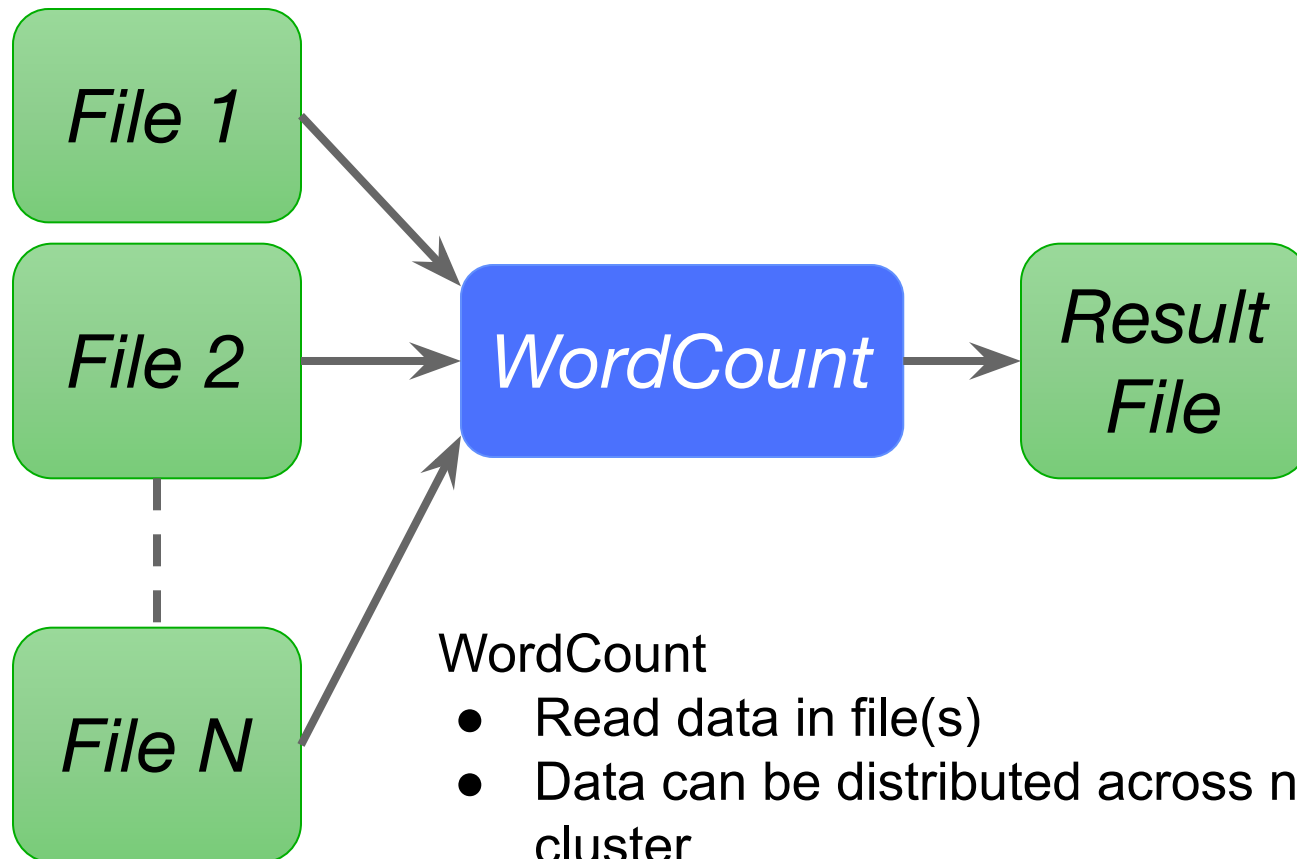
- DataFrames can be processed using 2 types of operations
 - **Transformation:** Creates new DataFrame from existing DataFrame
 - **Action:** Runs computation(s) on DataFrame and returns value



LAZY EVALUATION

- Transformations on DataFrames have **lazy evaluation**
 - Transformations are not immediately processed
 - Plan of operations is built
- Operations executed when **action** is performed
 - i.e., actions force computation
- Allows for optimizations in generating physical plan
- Example:
 - `filtered = strings.filter(strings["value"].contains("Spark"))`
 - Nothing is returned
 - `filtered.count()`
 - 'filter' is performed, and count is returned

WordCount

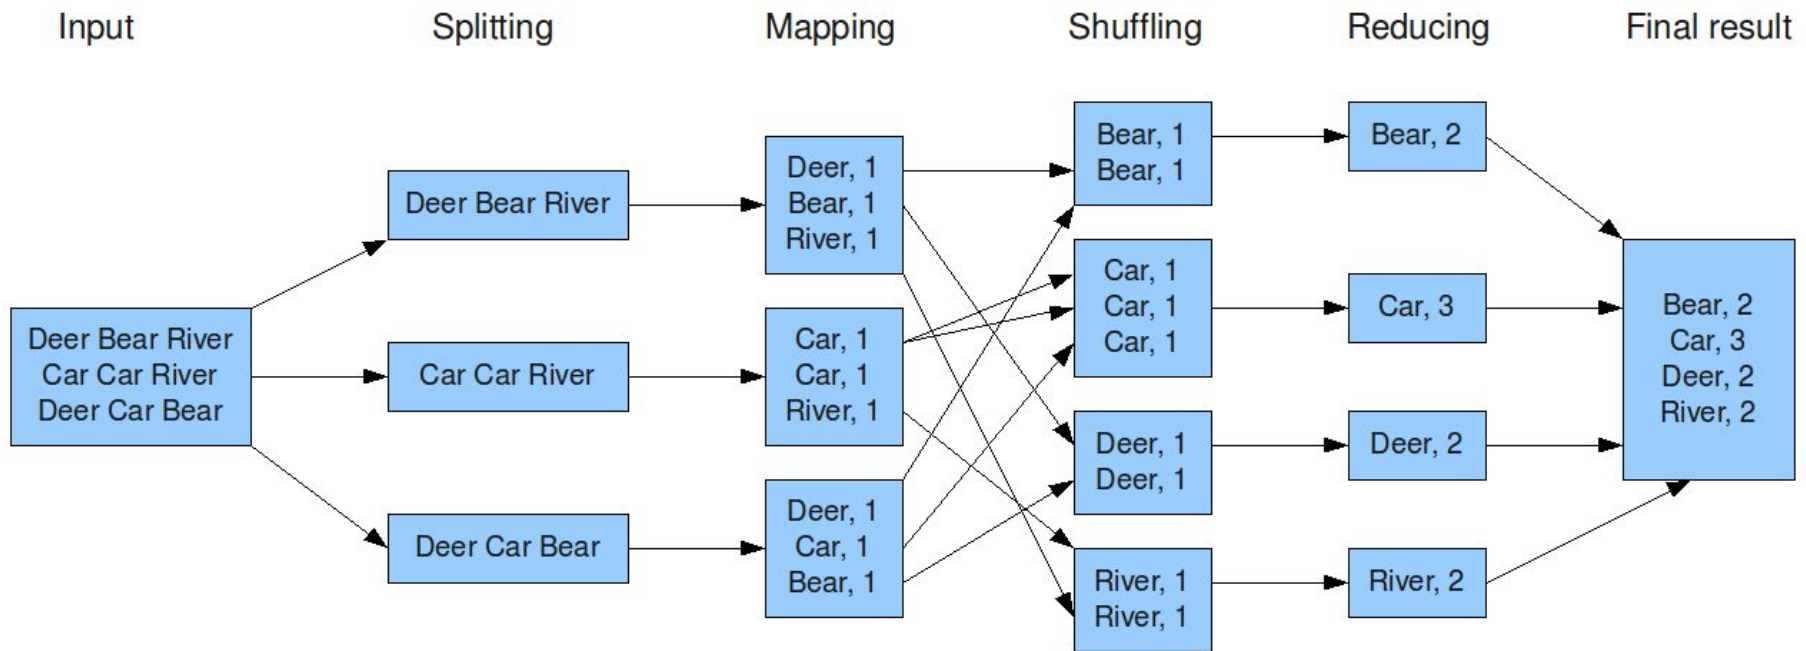


WordCount

- Read data in file(s)
- Data can be distributed across nodes in cluster
- Count number of occurrences of each word

WordCount

The overall MapReduce word count process



<https://www.todaysoftmag.com/article/1358/hadoop-mapreduce-deep-diving-and-tuning>

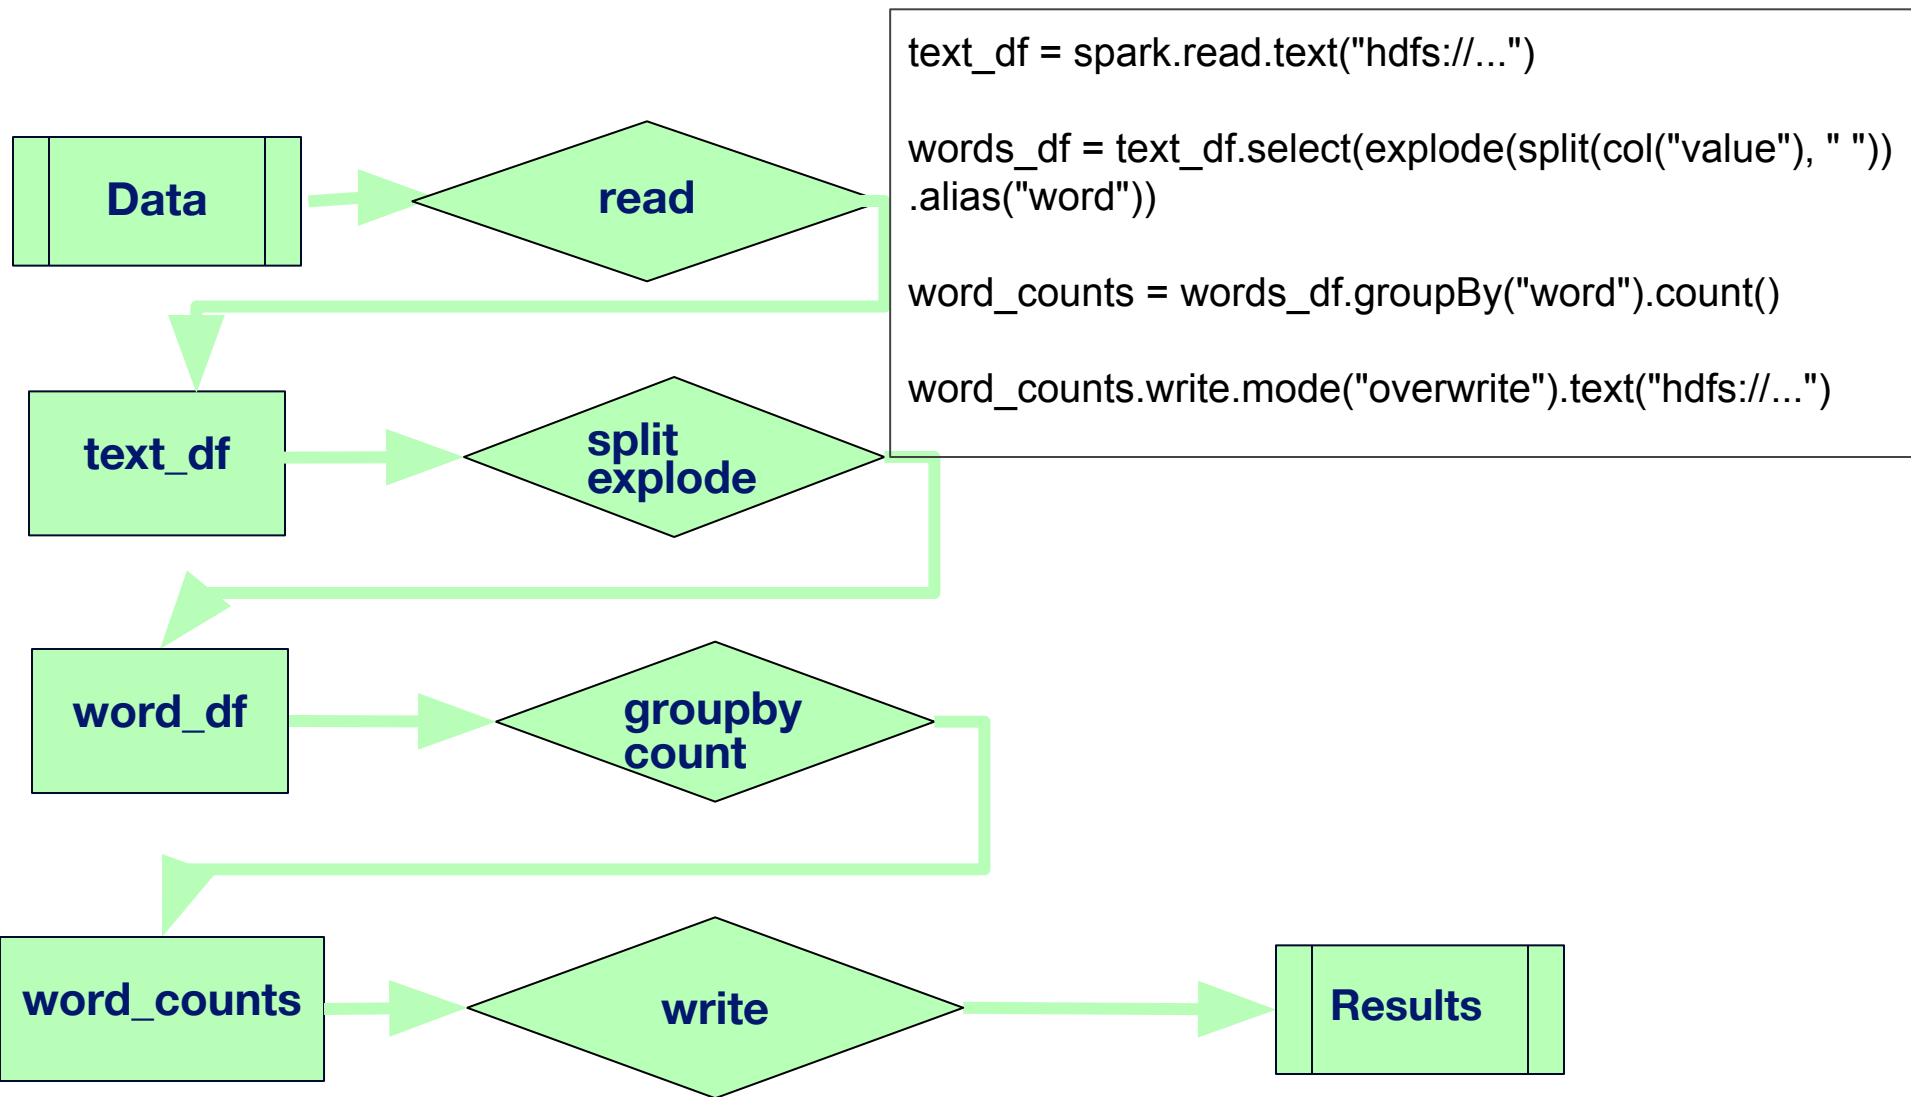
Data is split into
partitions

Map generates
key-value pairs

Pairs with same
key moved to
same partition

Reduce sums
values for
each key


WordCount



SPARK PROGRAM STRUCTURE


- **Start Spark session**
 - `spark = SparkSession.builder.config(conf=conf).getOrCreate()`
- **Create distributed dataset**
 - `df = spark.read.csv("data.csv",header="True")`
- **Apply transformations**
 - `new_df = df.filter(col("dept") == "Sales")`
- **Perform actions**
 - `df.collect()`
- **Stop Spark session**
 - `spark.stop()`

SPARK ARCHITECTURE



SPARK INTERFACE

Goals: speed, **ease of use**, generality, unified platform



WORDCOUNT EXAMPLE IN SPARK

Spark API available in Python, Scala, Java, and R

PySpark

```
text_df = spark.read.text("hdfs://...")
words_df = text_df.select(explode(split(col("value"), " ")).alias("word"))
word_counts = words_df.groupBy("word").count()
word_counts.write.mode("overwrite").text("hdfs://...")
```

SparkR

```
textDF <- read.text(spark, "hdfs://...")
wordsDF <- selectExpr(textDF, "explode(split(value, ' ')) as word")
wordCounts <- count(groupBy(wordsDF, "word"))
write.df(wordCounts, "hdfs://...", "text", mode = "overwrite")
```

Scala


```
val textDF = spark.read.text("hdfs://...")
val wordsDF = textDF.select(explode(split(col("value"), ""))).alias("word"))
val wordCounts = wordsDF.groupBy("word").count()
wordCounts.write.mode("overwrite").text("hdfs://...")
```

SPARK - GENERALITY

- Goals: speed, ease of use, **generality**, unified platform
- Support for several data sources
 - Local file systems, HDFS, RDBMSs, MongoDB, Kafka, AWS S3, etc.
- Can run on various platforms
 - Hadoop, Kubernetes, cloud, standalone
- Support for multiple workloads
 - batch, streaming
 - machine learning, SQL, graph processing


SPARK - UNIFIED PLATFORM

- Goals: speed, ease of use, generality, **unified platform**



- Provides unified platform for various analytics processing
- **Spark engine** provides core capabilities for scalable processing
- **Spark libraries** provide additional higher-level functionality for diverse workloads

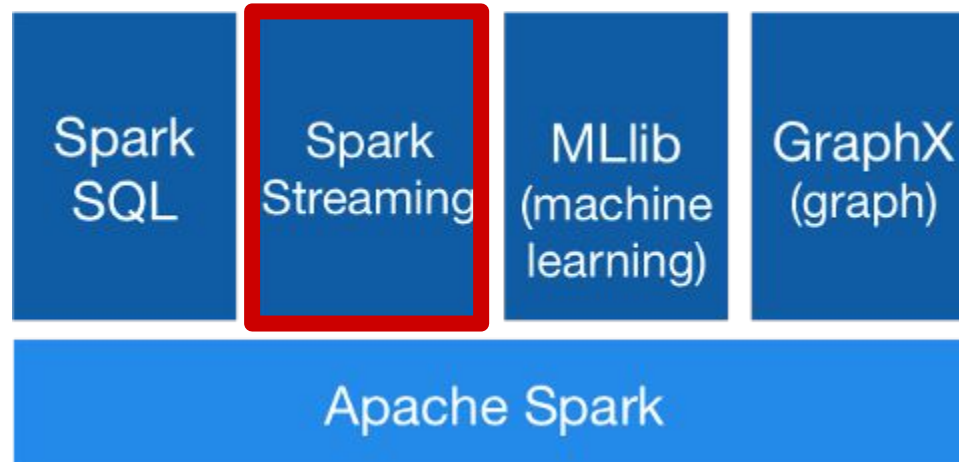
SPARK SQL



- **Structured Data Processing**

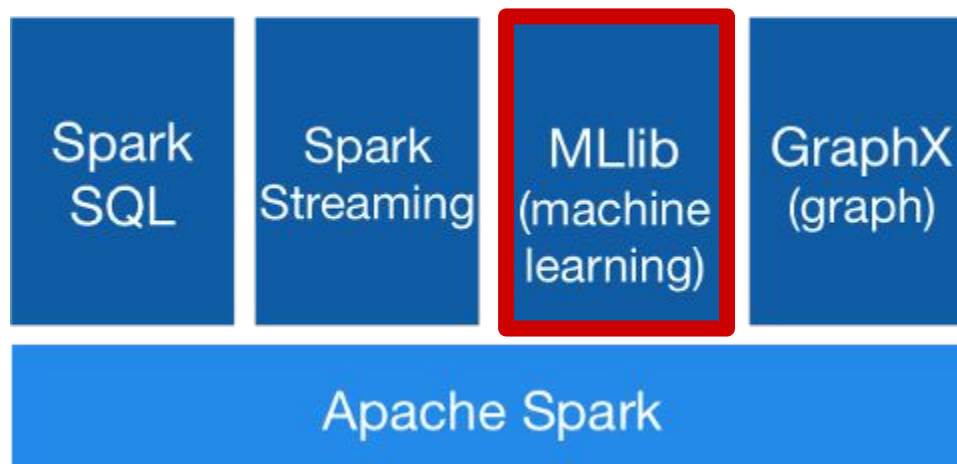
- Provides support for SQL and query processing
- Has APIs for SQL, Scala, Java, Python, and R
- Generated underlying code is identical

SPARK STREAMING



- **Streaming Data Processing**
 - Scalable processing for real-time analytics
 - Structured streaming
 - Data stream is divided into micro-batches of data
 - Same operations for static data can be used
 - Has APIs for Scala, Java, and Python

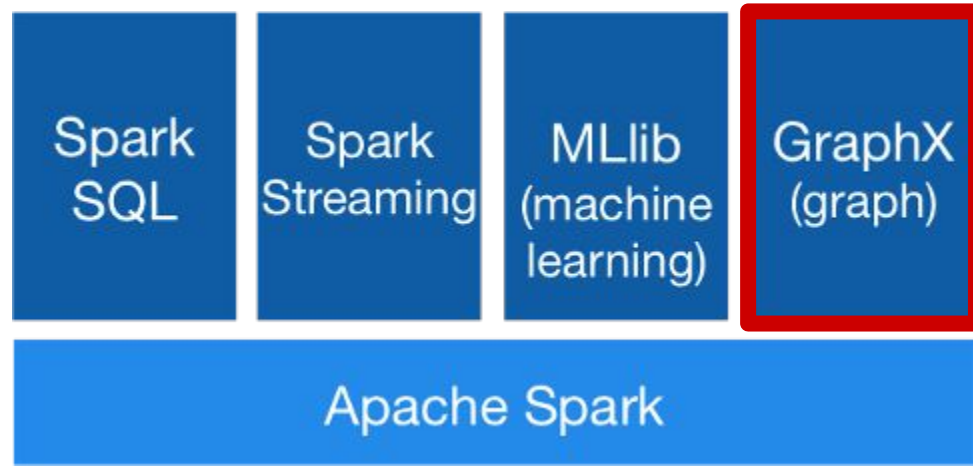
SPARK MLLIB



- **Machine Learning**

- Scalable machine learning library
- Scalable implementations of machine learning algorithms and utilities
- Has APIs for Scala, Java, Python, and R

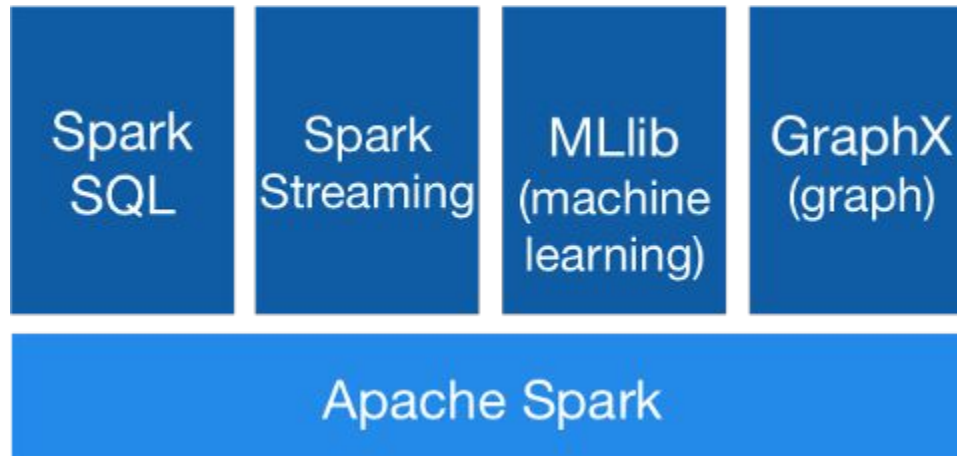
SPARK GRAPHX / GRAPHFRAMES



- **Graph Computation**

- Scalable graph processing
- Special structures for storing vertex and edge information & operations for manipulating graphs
- GraphX (RDD-based) & GraphFrames (DF-based)
- Has APIs in Scala, Java, Python (GraphFrames)

SPARK



Unified engine for large-scale data analytics

Goals: speed, ease of use, generality, unified platform

Spark Resources

- PySpark SQL Basics Cheat Sheet
 - PDF
- Spark Main Page
 - <https://spark.apache.org/>
- Spark Overview
 - <https://spark.apache.org/docs/latest/index.html>
- Spark Examples
 - <https://spark.apache.org/examples.html>
- Spark SQL, DataFrames and DataSets Programming Guide
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- Spark MLlib Programming Guide
 - <https://spark.apache.org/docs/latest/ml-guide.html>
- PySpark API Documentation
 - <https://spark.apache.org/docs/latest/api/python/index.html>

Spark Demo

Setup

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu -l <account>`

- **Pull latest from repo**

- `git clone` or `git pull`
- URL:
<https://github.com/sdsc-hpc-training-org/advanced-computing-webinars.git>
- Code location: Jan-14-2025-Scalable-ML/Spark

Server Setup for PySpark - Command Line

- **In terminal window**
 - `export PATH="/cm/shared/apps/sdsc/galileo:${PATH}"`
 - `jupyter-shared-spark`
 - Alias for: `galileo launch --account ${HPC_ACCOUNT} --reservation ${HPC_RESERVATION_CPU} --partition shared --cpus 4 --memory 16 --time-limit 02:00:00 --env-modules singularitypro --sif /cm/shared/apps/containers/singularity/spark/spark-latest.sif --bind /expance,/scratch,/cm --quiet`
- **To check queue**
 - `squeue -u $USER`

PySpark Scaling Hands-On

- **Data**
 - BookReviews_5M.txt
 - Source : <https://jmcauley.ucsd.edu/data/amazon/>
- **Notebook**
 - pyspark_demo.ipynb
- **To do**
 - Change number of cores: 1, 2, 4
 - Note difference in execution times
 - Run each configuration 3 times

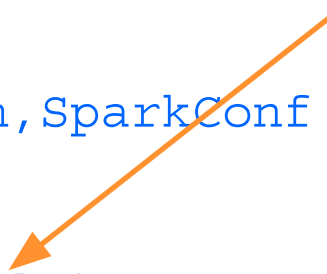
SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession, SparkConf

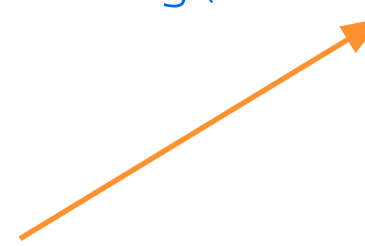
conf = SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])

spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

Use * to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one



GETTING EXECUTION TIMES

- In notebook, execution time is printed out in cell before Spark session is stopped (next to last cell)
- Need to restart the kernel and run all cells without stopping to get accurate execution time:
 - Run -> Restart Kernel and Run All Cells
- Find mean and standard deviation of execution times over 3 runs for
 - 1 core, 2 cores, and 4 cores

```
import pyspark
from pyspark.sql import SparkSession
```

```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[2]'),
    ('spark.app.name', 'PySpark Demo)])
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

Specify number of cores.
“*” uses all available cores



PySpark Cluster Analysis Hands-On

- **Data**

- Weather station measurements

- **Task**

- Perform cluster analysis to identify different weather patterns

- **Approach**

- Spark k-means

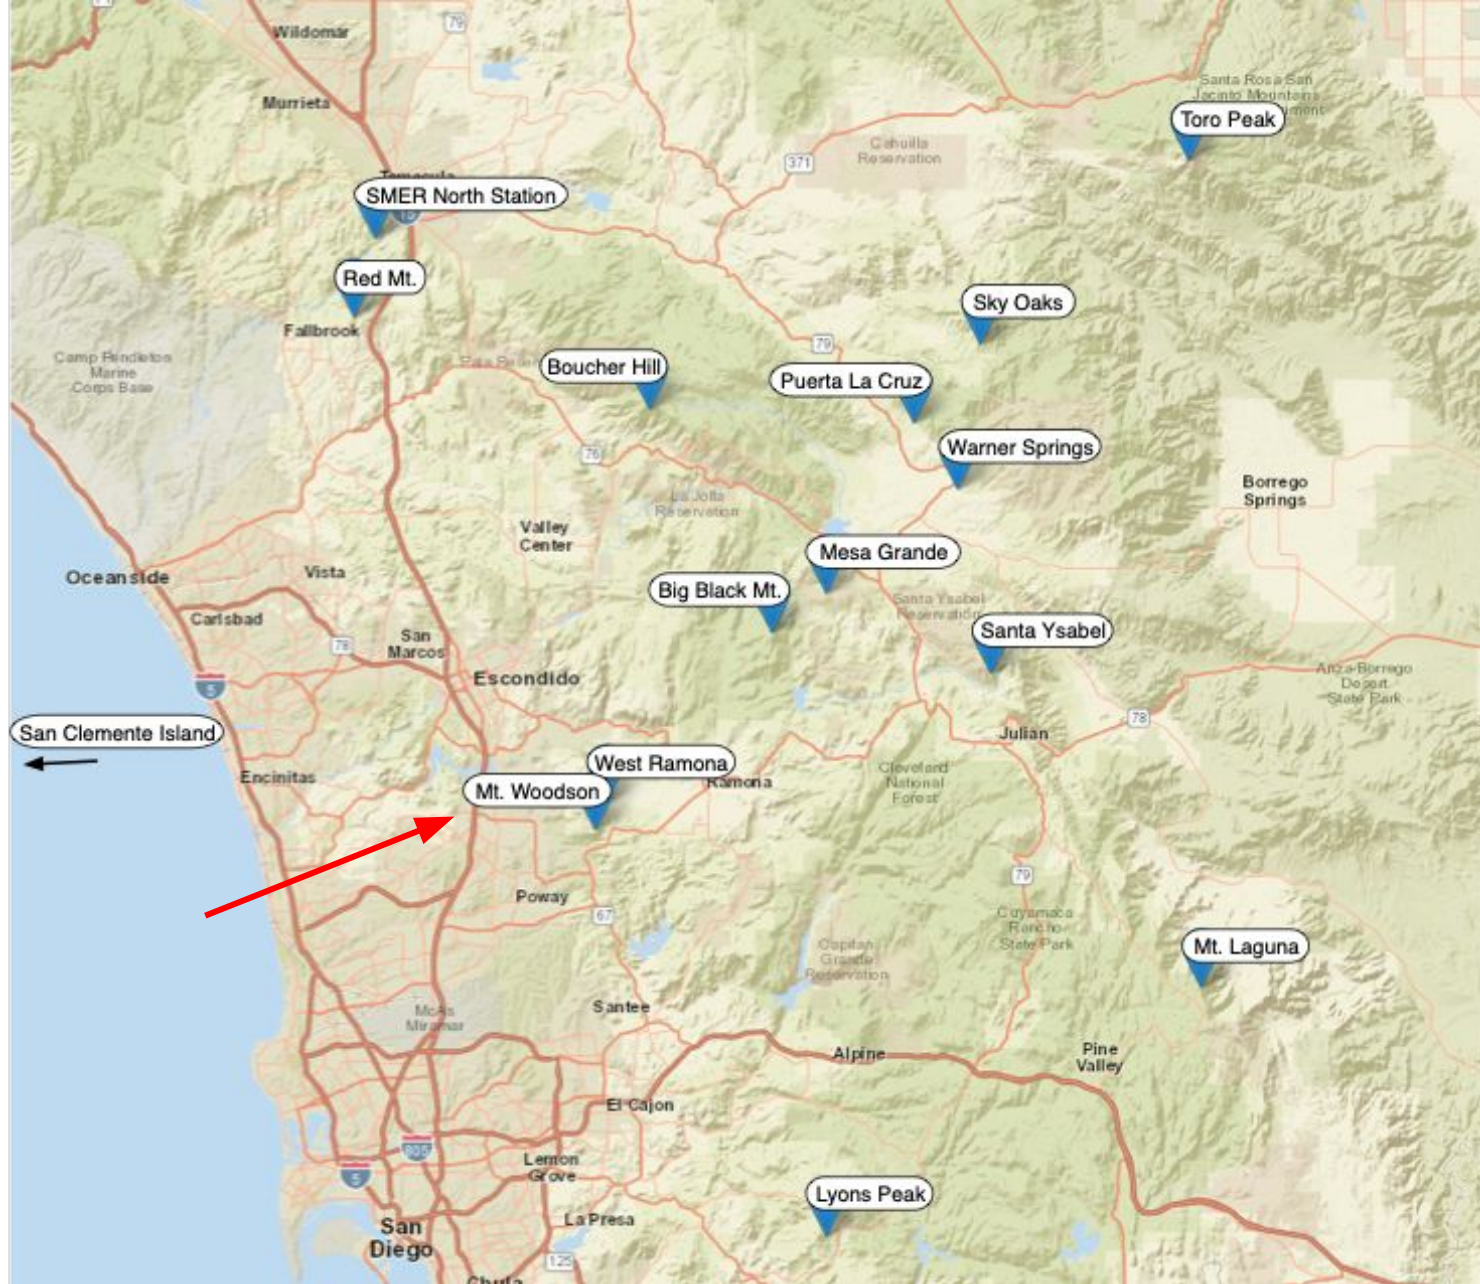
- **Notebooks**

- pyspark-clustering.ipynb # Starter notebook
- pyspark-clustering-wOutput.ipynb # Has cell outputs
- pyspark-clustering-soln.ipynb # Solution

Dataset Description

- **Measurements from weather station on Mt. Woodson, San Diego**
- **Air temperature, humidity, wind speed, wind direction, etc.**
- **Three years of data: Sep. 2011 - Sep. 2014**
 - minute_weather.csv: measurement every minute
- **Source**
 - <http://hpwren.ucsd.edu>

Map of HPWREN Weather Stations



Clustering Hands-On Overview

- **Setup**
 - Start Spark
 - Load modules
- **Load data**
 - Specify schema
 - Read in data from “minute_weather.csv”
- **Explore data**
 - Look at schema, number of rows, summary statistics
- **Prepare data**
 - Drop nulls
 - Create feature vector
- **Perform k-means cluster analysis**
 - Use elbow plot to determine k
 - Build k-means model
- **Evaluate clusters**
 - Plot cluster profiles
- **Stop Spark session**

Resources

- **Spark**
 - <https://spark.apache.org/>
- **PySpark API**
 - <https://spark.apache.org/docs/latest/api/python/index.html>
- **Spark DataFrame**
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- **MLlib**
 - <https://spark.apache.org/mllib/>
- **User's Guide**
 - https://spark.apache.org/docs/latest/api/python/user_guide/pandas_on_spark/index.html